

# Speaker Recognition

Term project for Machine Learning UoT course

The objective of this Notebook is to predict the speaker of an English digit using a Keras model.

This Notebook has three main sections.

1. Extract the features of the WAV files, save them into CSV files and store them into Pandas
2. The model will be trained based on the data set created by <https://github.com/Jakobovski/free-spoken-digit-dataset> (<https://github.com/Jakobovski/free-spoken-digit-dataset>)
3. The model will be re-trained based on the previous data plus the recordings made by Ankor (Indian accent), Caroline (Canadian female child accent) and Rodolfo (Brazilian accent)

In [ ]:

## Section 1

The output of this section is the CSV files with the data to be handle by the model

```
trainData      : ../data/recordings/train
testData       : ../data/recordings/test
moreTrainData  : ../data/recordings/moreSpeakersTrain
moreTestData   : ../data/recordings/moreSpeakersTest
```

```
In [1]: # If true, the WAV files will be read and their features will be saved in the
        CSV files
        # As this is the most time consuming task, only enable it if you don't have th
        e CSV files yet
        CREATE_CSV_FILES = True
```

```
In [2]: # Defines the names of the CSV files
        TRAIN_CSV_FILE = "train.csv"
        TEST_CSV_FILE = "test.csv"
        MORE_TRAIN_CSV_FILE = "more_train.csv"
        MORE_TEST_CSV_FILE = "more_test.csv"
```

```

In [3]: import matplotlib.pyplot as plt
import numpy as np
from matplotlib import cm
import librosa
import csv
import os

def extractWavFeatures(soundFilesFolder, csvFileName):
    print("The features of the files in the folder "+soundFilesFolder+" will be saved to "+csvFileName)
    header = 'filename chroma_stft rmse spectral_centroid spectral_bandwidth rolloff zero_crossing_rate'
    for i in range(1, 21):
        header += f' mfcc{i}'
    header += ' label'
    header = header.split()
    print('CSV Header: ', header)
    file = open(csvFileName, 'w', newline='')
    #with file:
    writer = csv.writer(file)
    writer.writerow(header)
    genres = '1 2 3 4 5 6 7 8 9 0'.split()
    for filename in os.listdir(soundFilesFolder):
        number = f'{soundFilesFolder}/{filename}'
        y, sr = librosa.load(number, mono=True, duration=30)
        # remove leading and trailing silence
        y, index = librosa.effects.trim(y)
        chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
        rmse = librosa.feature.rms(y=y)
        spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
        spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
        rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
        zcr = librosa.feature.zero_crossing_rate(y)
        mfcc = librosa.feature.mfcc(y=y, sr=sr)
        to_append = f'{filename} {np.mean(chroma_stft)} {np.mean(rmse)} {np.mean(spec_cent)} {np.mean(spec_bw)} {np.mean(rolloff)} {np.mean(zcr)}'
        for e in mfcc:
            to_append += f' {np.mean(e)}'
        writer.writerow(to_append.split())
    file.close()
    print("End of extractWavFeatures")

if (CREATE_CSV_FILES == True):
    extractWavFeatures("../data/recordings/train", TRAIN_CSV_FILE)
    extractWavFeatures("../data/recordings/test", TEST_CSV_FILE)
    extractWavFeatures("../data/recordings/moreSpeakersTrain", MORE_TRAIN_CSV_FILE)
    extractWavFeatures("../data/recordings/moreSpeakersTest", MORE_TEST_CSV_FILE)
    print("CSV files are created")
else:
    print("CSV files creation is skipped")

```

The features of the files in the folder ../data/recordings/train will be saved to train.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/test will be saved to test.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/moreSpeakersTrain will be saved to more\_train.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/moreSpeakersTest will be saved to more\_test.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

CSV files are created

In [4]: *#Reading a dataset and convert file name to corresponding umnber*

```
import pandas as pd
import csv
from sklearn import preprocessing

def preProcessData(csvFileName):
    print(csvFileName+ " will be preprocessed")
    data = pd.read_csv(csvFileName)
    # we have six speakers:
    # 0: Jackson
    # 1: Nicolas
    # 2: Theo
    # 3: Ankur
    # 4: Caroline
    # 5: Rodolfo
    filenameArray = data['filename']
    speakerArray = []
    #print(filenameArray)
    for i in range(len(filenameArray)):
        speaker = filenameArray[i][2]
        #print(speaker)
        if speaker == "j":
            speaker = "0"
        elif speaker == "n":
            speaker = "1"
        elif speaker == "t":
            speaker = "2"
        elif speaker == "a":
            speaker = "3"
        elif speaker == "c":
            speaker = "4"
        elif speaker == "r":
            speaker = "5"
        else:
            speaker = "6"
        #print(speaker)
        speakerArray.append(speaker)
    data['number'] = speakerArray
    #Dropping unnecessary columns
    data = data.drop(['filename'],axis=1)
    data = data.drop(['label'],axis=1)
    data = data.drop(['chroma_stft'],axis=1)
    data.shape

    print("Preprocessing is finished")
    print(data.head())
    return data

trainData = preProcessData(TRAIN_CSV_FILE)
testData = preProcessData(TEST_CSV_FILE)
moreTrainData = preProcessData(MORE_TRAIN_CSV_FILE)
moreTestData = preProcessData(MORE_TEST_CSV_FILE)
```

train.csv will be preprocessed

Preprocessing is finished

	rmse	spectral_centroid	spectral_bandwidth	rolloff	\
0	0.112672	741.829081	758.492178	1438.494873	
1	0.090344	635.610880	670.336296	1160.452403	
2	0.091456	667.786694	732.606545	1257.180176	
3	0.087751	712.304185	731.292437	1449.104818	
4	0.096603	844.363886	777.868127	1569.583263	

	zero_crossing_rate	mfcc1	mfcc2	mfcc3	mfcc4	\
0	0.034023	-295.578461	189.853683	-19.606564	6.078507	
1	0.033458	-339.148743	204.005249	-7.485526	14.297898	
2	0.033268	-327.507416	195.596924	-3.994768	21.315840	
3	0.035916	-320.809937	200.023743	-8.186146	12.661074	
4	0.049465	-315.801300	195.674118	-13.324564	3.544238	

	mfcc5	...	mfcc12	mfcc13	mfcc14	mfcc15	mfcc16	\
0	22.067095	...	-25.725817	-5.172223	-8.323026	-10.299589	-0.144793	
1	20.885136	...	-23.196365	-1.290891	-5.515564	-15.416287	0.405876	
2	18.372593	...	-18.677113	-3.098450	-10.447586	-10.053793	3.248016	
3	15.654718	...	-20.832333	-1.118007	-6.681235	-11.685319	2.010999	
4	12.279986	...	-18.158249	6.031695	-6.353736	-15.983871	1.465030	

	mfcc17	mfcc18	mfcc19	mfcc20	number
0	-9.017329	-4.569392	2.881349	-15.627436	0
1	-3.624587	-11.204143	-0.096359	-6.751650	0
2	-11.686995	-10.726046	6.857377	-9.067446	0
3	-5.946658	-6.905020	4.136240	-9.614882	0
4	-5.109472	-8.666434	5.026890	-5.346444	0

[5 rows x 26 columns]

test.csv will be preprocessed

Preprocessing is finished

	rmse	spectral_centroid	spectral_bandwidth	rolloff	\
0	0.095394	756.450712	761.875940	1463.941148	
1	0.040176	791.046914	1039.695939	2027.709961	
2	0.006984	958.934867	941.639039	2084.106445	
3	0.071547	759.877794	899.957003	1427.553489	
4	0.030382	968.793389	1024.834851	1911.968994	

	zero_crossing_rate	mfcc1	mfcc2	mfcc3	mfcc4	\
0	0.037296	-328.263885	180.479416	-0.485355	15.525293	
1	0.031440	-385.602570	189.328186	-37.268154	59.937920	
2	0.040039	-542.812622	217.971329	-62.197266	21.537390	
3	0.030429	-355.530396	204.388977	-20.676432	26.671131	
4	0.045654	-376.499390	237.137833	-59.964413	37.715607	

	mfcc5	...	mfcc12	mfcc13	mfcc14	mfcc15	mfcc16	\
0	20.992447	...	-15.426966	7.284101	-6.443027	-13.377846	-2.407696	
1	45.049831	...	-31.051588	3.420474	-9.762264	-11.220519	12.306476	
2	37.756233	...	-26.797358	5.341060	-12.159102	-14.180812	9.346475	
3	15.797892	...	-18.198524	4.029843	-10.552087	-21.039103	-5.634320	
4	21.510382	...	-21.805593	13.740063	-6.738161	-9.305484	13.205662	

	mfcc17	mfcc18	mfcc19	mfcc20	number
0	-12.902534	-10.437113	-1.025342	-15.457672	0
1	-5.082399	-3.775387	9.707520	-8.757109	1

```

2 -10.899978 -9.715154 5.997578 -12.574761 2
3 -14.788965 -11.016036 -1.313916 -16.993853 0
4 -11.917943 -7.877903 9.777577 -10.397771 1

```

[5 rows x 26 columns]

more\_train.csv will be preprocessed

Preprocessing is finished

```

      rmse  spectral_centroid  spectral_bandwidth  rolloff \
0  0.039759      1358.208628      1890.243941  2540.917969
1  0.302424       879.994019      1137.986581  1826.501859
2  0.026959      1237.544903      1219.890113  2372.379244
3  0.027274      1523.814892      2030.693021  3053.946533
4  0.304633       625.579402       799.806332  1029.825439

```

```

      zero_crossing_rate      mfcc1      mfcc2      mfcc3      mfcc4 \
0      0.093363 -349.631744  135.204880  19.397516  12.044560
1      0.031423 -225.136642  150.581146 -11.930015   5.277394
2      0.059871 -378.096527  180.817047 -41.006123  22.502394
3      0.084253 -403.760406  127.670433  23.293980   7.495267
4      0.023584 -228.231903  191.904144   1.268919  17.871376

```

```

      mfcc5 ...      mfcc12      mfcc13      mfcc14      mfcc15      mfcc16 \
0 -8.521679 ... -12.872719 -3.118359 -4.095297 -8.339793 -4.422189
1 -16.297689 ... -13.114679 -13.413434 -17.883400 -17.694012 -18.091360
2 -8.963315 ... -32.601864 -4.715013 -8.889856 -11.225335 -10.951420
3 -6.774144 ... -6.803737 -3.132436 -2.999972 -14.086405 -9.653265
4 -4.571987 ... -5.878148 -9.646009 -15.565687 -17.903820 -14.293053

```

```

      mfcc17      mfcc18      mfcc19      mfcc20  number
0 -2.988979 -0.864654 -4.008632  3.243911      3
1 -14.684992 -16.672678 -10.986701 -10.445865      4
2 -8.972343 -4.716431 -5.655877 -7.133625      5
3 -7.696736 -2.747038 -11.928693 -4.345026      3
4 -14.907639 -9.752651 -10.333107 -7.938378      4

```

[5 rows x 26 columns]

more\_test.csv will be preprocessed

Preprocessing is finished

```

      rmse  spectral_centroid  spectral_bandwidth  rolloff \
0  0.070496      1736.761057      1697.429353  2782.932447
1  0.178100       986.240750      1243.094994  2207.479581
2  0.159381      1140.769641      1151.467042  2283.347731
3  0.102912      1270.788895      1624.572386  2442.362154
4  0.243922       839.268659       985.135107  1423.724724

```

```

      zero_crossing_rate      mfcc1      mfcc2      mfcc3      mfcc4 \
0      0.095979 -270.462311  120.525047 -16.563780  25.340765
1      0.045351 -230.533142  171.157135 -5.515658  20.675375
2      0.056547 -261.164734  169.949051 -12.116495   0.249355
3      0.055965 -264.190643  154.578873 -9.901300   1.276820
4      0.037626 -191.595886  189.683578 -29.157555  16.972124

```

```

      mfcc5 ...      mfcc12      mfcc13      mfcc14      mfcc15      mfcc16 \
0 -8.705835 ... -20.140141 -1.016966 -11.883263 -8.389129 -4.095136
1 -3.108402 ... -14.097899 -9.547957 -9.970568 -15.526528 -17.028362
2  8.262510 ... -26.140615 -3.177802 -5.860640 -14.931437 -12.212286
3 -4.987144 ... -9.880266 -0.010168 -10.896443 -5.682155 -0.433506

```

```
4 -13.665516 ... -14.553379 -6.711682 -15.045198 -11.206120 -11.227142
```

	mfcc17	mfcc18	mfcc19	mfcc20	number
0	2.231521	-6.269879	-7.276823	0.429325	3
1	-13.014733	-9.094010	1.289093	-5.188507	4
2	-5.934423	-2.185390	-7.319934	-12.082070	5
3	6.786911	-10.947194	-11.484864	-0.297753	3
4	-9.380514	-11.611271	-7.099937	-6.555156	4

```
[5 rows x 26 columns]
```

In [ ]:

## Section 2

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings)

Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings for each digit for each speaker: 30 recordings total.

The data used here comes from the recordings stored in:

- ../data/recordings/train
- ../data/recordings/test

The model will be trained to predict the speaker of a digit.

```
In [5]: # Splitting the dataset into training, validation and testing dataset
from sklearn.model_selection import train_test_split
X = np.array(trainData.iloc[:, :-1], dtype = float)
y = trainData.iloc[:, -1]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_
state=42)

X_test = np.array(testData.iloc[:, :-1], dtype = float)
y_test = testData.iloc[:, -1]

print("Y from training data:", y_train.shape)
print("Y from validation data:", y_val.shape)
print("Y from test data:", y_test.shape)

Y from training data: (1029,)
Y from validation data: (441,)
Y from test data: (30,)
```

```
In [6]: #Normalizing the dataset
from sklearn.preprocessing import StandardScaler
import numpy as np
scaler = StandardScaler()
X_train = scaler.fit_transform( X_train )
X_val = scaler.transform( X_val )
X_test = scaler.transform( X_test )

print("X from training data", X_train.shape)
print("X from validation data", X_val.shape)
print("X from test data", X_test.shape)
```

```
X from training data (1029, 25)
X from validation data (441, 25)
X from test data (30, 25)
```



```
In [7]: #Creating a Model
from keras import models
from keras import layers
import keras

# model 1
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_shape=(X_train.shape[1],)))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# Learning Process of a model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# simple early stopping
from keras.callbacks import EarlyStopping

es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)

#Train with early stopping to avoid overfitting
history = model.fit(X_train,
                    y_train,
                    validation_data=(X_val, y_val),
                    epochs=50,
                    batch_size=128,
                    callbacks=[es])
```

Using TensorFlow backend.

Train on 1029 samples, validate on 441 samples

Epoch 1/50

1029/1029 [=====] - 0s 187us/step - loss: 2.1741 - accuracy: 0.2439 - val\_loss: 1.5546 - val\_accuracy: 0.7029

Epoch 2/50

1029/1029 [=====] - 0s 31us/step - loss: 1.4414 - accuracy: 0.5384 - val\_loss: 0.8647 - val\_accuracy: 0.9274

Epoch 3/50

1029/1029 [=====] - 0s 33us/step - loss: 0.9387 - accuracy: 0.6774 - val\_loss: 0.3857 - val\_accuracy: 0.9728

Epoch 4/50

1029/1029 [=====] - 0s 37us/step - loss: 0.5667 - accuracy: 0.8066 - val\_loss: 0.1855 - val\_accuracy: 0.9751

Epoch 5/50

1029/1029 [=====] - 0s 48us/step - loss: 0.4205 - accuracy: 0.8669 - val\_loss: 0.1118 - val\_accuracy: 0.9751

Epoch 6/50

1029/1029 [=====] - 0s 42us/step - loss: 0.2702 - accuracy: 0.9145 - val\_loss: 0.0831 - val\_accuracy: 0.9773

Epoch 7/50

1029/1029 [=====] - 0s 46us/step - loss: 0.2007 - accuracy: 0.9475 - val\_loss: 0.0588 - val\_accuracy: 0.9864

Epoch 8/50

1029/1029 [=====] - 0s 60us/step - loss: 0.1665 - accuracy: 0.9582 - val\_loss: 0.0507 - val\_accuracy: 0.9887

Epoch 9/50

1029/1029 [=====] - 0s 52us/step - loss: 0.1537 - accuracy: 0.9572 - val\_loss: 0.0456 - val\_accuracy: 0.9909

Epoch 10/50

1029/1029 [=====] - 0s 56us/step - loss: 0.1288 - accuracy: 0.9611 - val\_loss: 0.0438 - val\_accuracy: 0.9909

Epoch 11/50

1029/1029 [=====] - 0s 86us/step - loss: 0.1228 - accuracy: 0.9640 - val\_loss: 0.0393 - val\_accuracy: 0.9909

Epoch 12/50

1029/1029 [=====] - 0s 56us/step - loss: 0.0626 - accuracy: 0.9864 - val\_loss: 0.0326 - val\_accuracy: 0.9887

Epoch 13/50


1029/1029 [=====] - 0s 53us/step - loss: 0.0822 - accuracy: 0.9757 - val\_loss: 0.0300 - val\_accuracy: 0.9887

Epoch 14/50

1029/1029 [=====] - 0s 49us/step - loss: 0.0931 - accuracy: 0.9708 - val\_loss: 0.0316 - val\_accuracy: 0.9909

Epoch 00014: early stopping

```
In [8]: # plot training history
from matplotlib import pyplot
pyplot.plot(history.history['loss'], label='train')
pyplot.plot(history.history['val_loss'], label='test')
pyplot.legend()
pyplot.show()
```



## Auxiliary functions to show the results

```
In [9]: def getSpeaker(speaker):
    speaker = str(speaker)
    if speaker == "0":
        return "Jackson"
    elif speaker == "1":
        return "Nicola"
    elif speaker == "2":
        return "Theo"
    elif speaker == "3":
        return "Ankur"
    elif speaker == "4":
        return "Caroline"
    elif speaker == "5":
        return "Rodolfo"
    else:
        speaker = "Unknown"

def printPrediction(X_data, y_data, printDigit):
    print('\n# Generate predictions')
    for i in range(len(y_data)):
        prediction = getSpeaker(model.predict_classes(X_data[i:i+1])[0])
        speaker = getSpeaker(y_data[i])
        if printDigit == True:
            print("Number={0:d}, y={1:10s}- prediction={2:10s}- match={3}".format(i, speaker, prediction, speaker==prediction))
        else:
            print("y={0:10s}- prediction={1:10s}- match={2}".format(speaker, prediction, speaker==prediction))
```

```
In [10]: import numpy as np
from keras import backend as K
from keras.models import Sequential
from keras.layers.core import Dense, Dropout, Activation, Flatten
from keras.layers.convolutional import Convolution2D, MaxPooling2D
from keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import classification_report, confusion_matrix

def report(X_data, y_data):
    #Confution Matrix and Classification Report
    Y_pred = model.predict_classes(X_data)
    y_test_num = y_data.astype(np.int64)
    conf_mt = confusion_matrix(y_test_num, Y_pred)
    print(conf_mt)
    plt.matshow(conf_mt)
    plt.show()
    print('\nClassification Report')
    target_names = ["Jackson", "Nicola", "Theo", "Ankur", "Caroline", "Rodolfo", "Unknown"]
    print(classification_report(y_test_num, Y_pred))
```

```
In [ ]:
```

## Present the model performance

```
In [11]: print('\n# TEST DATA #\n')
score = model.evaluate(X_test, y_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_test[0:10], y_test[0:10], False)

# TEST DATA #


30/30 [=====] - 0s 33us/step
accuracy: 100.00%

# Generate predictions
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
```

```
In [12]: print("Classification Report for Test Data\n")
report(X_test, y_test)
```

Classification Report for Test Data

```
[[10  0  0]
 [ 0 10  0]
 [ 0  0 10]]
```



Classification Report					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	10	
1	1.00	1.00	1.00	10	
2	1.00	1.00	1.00	10	
accuracy			1.00	30	
macro avg	1.00	1.00	1.00	30	
weighted avg	1.00	1.00	1.00	30	

```
In [ ]:
```

## Section 3

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings)  
Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings for each digit for each speaker: 30 recordings total.

In addition, there are 2 recordings for each digit for each speaker: Ankur, Caroline and Rodolfo (total 60 recordings) This addition training data has 1 recordings for each digit for each speaker: 30 recordings total. This addition test data has 1 recordings for each digit for each speaker: 30 recordings total.

Therefore the full data set has:

- Training: 1500 recordings
- Training: 60 recordings

The data used here comes from the recordings stored in:

- ../data/recordings/train
- ../data/recordings/test
- ../data/recordings/moreSpeakersTrain
- ../data/recordings/moreSpeakersTest

In [ ]:

```
In [13]: # Splitting the dataset into training, validation and testing dataset
from sklearn.model_selection import train_test_split

fullTrainData = trainData.append(moreTrainData)

X = np.array(fullTrainData.iloc[:, :-1], dtype = float)
y = fullTrainData.iloc[:, -1]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_
state=42)

X_test = np.array(testData.iloc[:, :-1], dtype = float)
y_test = testData.iloc[:, -1]

X_more_test = np.array(moreTestData.iloc[:, :-1], dtype = float)
y_more_test = moreTestData.iloc[:, -1]

print("Y from training data:", y_train.shape)
print("Y from validation data:", y_val.shape)
print("Y from test data:", y_test.shape)
print("Y from other speakers test data:", y_more_test.shape)

Y from training data: (1050,)
Y from validation data: (450,)
Y from test data: (30,)
Y from other speakers test data: (30,)
```

```
In [14]: #Normalizing the dataset
from sklearn.preprocessing import StandardScaler
import numpy as np
scaler = StandardScaler()
X_train = scaler.fit_transform( X_train )
X_val = scaler.transform( X_val )
X_test = scaler.transform( X_test )
X_more_test = scaler.transform( X_more_test )

print("X from training data", X_train.shape)
print("X from validation data", X_val.shape)
print("X from test data", X_test.shape)
print("X from other speakers test data", X_more_test.shape)
```

```
X from training data (1050, 25)
X from validation data (450, 25)
X from test data (30, 25)
X from other speakers test data (30, 25)
```

```
In [15]: #Creating a Model
from keras import models
from keras import layers
import keras

# model 1
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_shape=(X_train.shape[1],)))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# Learning Process of a model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# simple early stopping
from keras.callbacks import EarlyStopping

es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)

#Train with early stopping to avoid overfitting
history = model.fit(X_train,
                    y_train,
                    validation_data=(X_val, y_val),
                    epochs=50,
                    batch_size=128,
                    callbacks=[es])
```



Train on 1050 samples, validate on 450 samples

Epoch 1/50

1050/1050 [=====] - 0s 176us/step - loss: 2.0021 - accuracy: 0.2886 - val\_loss: 1.4814 - val\_accuracy: 0.8289

Epoch 2/50

1050/1050 [=====] - 0s 30us/step - loss: 1.3657 - accuracy: 0.5495 - val\_loss: 0.9358 - val\_accuracy: 0.9111

Epoch 3/50

1050/1050 [=====] - 0s 30us/step - loss: 0.9921 - accuracy: 0.6943 - val\_loss: 0.6088 - val\_accuracy: 0.9378

Epoch 4/50

1050/1050 [=====] - 0s 31us/step - loss: 0.7191 - accuracy: 0.8000 - val\_loss: 0.3982 - val\_accuracy: 0.9400

Epoch 5/50

1050/1050 [=====] - 0s 36us/step - loss: 0.5295 - accuracy: 0.8667 - val\_loss: 0.2920 - val\_accuracy: 0.9489

Epoch 6/50

1050/1050 [=====] - 0s 32us/step - loss: 0.3837 - accuracy: 0.9105 - val\_loss: 0.2339 - val\_accuracy: 0.9533

Epoch 7/50

1050/1050 [=====] - 0s 34us/step - loss: 0.3030 - accuracy: 0.9229 - val\_loss: 0.1947 - val\_accuracy: 0.9556

Epoch 8/50

1050/1050 [=====] - 0s 34us/step - loss: 0.2402 - accuracy: 0.9495 - val\_loss: 0.1559 - val\_accuracy: 0.9578

Epoch 9/50

1050/1050 [=====] - 0s 36us/step - loss: 0.2211 - accuracy: 0.9524 - val\_loss: 0.1291 - val\_accuracy: 0.9622

Epoch 10/50

1050/1050 [=====] - 0s 37us/step - loss: 0.1796 - accuracy: 0.9524 - val\_loss: 0.1190 - val\_accuracy: 0.9644

Epoch 11/50

1050/1050 [=====] - 0s 36us/step - loss: 0.1548 - accuracy: 0.9629 - val\_loss: 0.1083 - val\_accuracy: 0.9644

Epoch 12/50

1050/1050 [=====] - 0s 35us/step - loss: 0.1522 - accuracy: 0.9667 - val\_loss: 0.0997 - val\_accuracy: 0.9689

Epoch 13/50

1050/1050 [=====] - 0s 36us/step - loss: 0.1574 - accuracy: 0.9590 - val\_loss: 0.0904 - val\_accuracy: 0.9733

Epoch 14/50

1050/1050 [=====] - 0s 35us/step - loss: 0.1350 - accuracy: 0.9667 - val\_loss: 0.0844 - val\_accuracy: 0.9733

Epoch 15/50

1050/1050 [=====] - 0s 37us/step - loss: 0.1133 - accuracy: 0.9733 - val\_loss: 0.0812 - val\_accuracy: 0.9733

Epoch 16/50

1050/1050 [=====] - 0s 36us/step - loss: 0.1086 - accuracy: 0.9714 - val\_loss: 0.0767 - val\_accuracy: 0.9733

Epoch 17/50

1050/1050 [=====] - 0s 34us/step - loss: 0.1026 - accuracy: 0.9762 - val\_loss: 0.0718 - val\_accuracy: 0.9756

Epoch 18/50

1050/1050 [=====] - 0s 35us/step - loss: 0.0946 - accuracy: 0.9695 - val\_loss: 0.0697 - val\_accuracy: 0.9756

Epoch 19/50

1050/1050 [=====] - 0s 37us/step - loss: 0.0793 - ac

```


curacy: 0.9705 - val_loss: 0.0678 - val_accuracy: 0.9756
Epoch 20/50
1050/1050 [=====] - 0s 38us/step - loss: 0.0859 - ac
curacy: 0.9743 - val_loss: 0.0621 - val_accuracy: 0.9800
Epoch 21/50
1050/1050 [=====] - 0s 33us/step - loss: 0.0705 - ac
curacy: 0.9752 - val_loss: 0.0591 - val_accuracy: 0.9800
Epoch 22/50
1050/1050 [=====] - 0s 37us/step - loss: 0.0796 - ac
curacy: 0.9771 - val_loss: 0.0572 - val_accuracy: 0.9800
Epoch 23/50
1050/1050 [=====] - 0s 35us/step - loss: 0.0750 - ac
curacy: 0.9733 - val_loss: 0.0522 - val_accuracy: 0.9822
Epoch 24/50
1050/1050 [=====] - 0s 35us/step - loss: 0.0571 - ac
curacy: 0.9838 - val_loss: 0.0496 - val_accuracy: 0.9822
Epoch 25/50
1050/1050 [=====] - 0s 38us/step - loss: 0.0567 - ac
curacy: 0.9800 - val_loss: 0.0478 - val_accuracy: 0.9844
Epoch 26/50
1050/1050 [=====] - 0s 40us/step - loss: 0.0463 - ac
curacy: 0.9857 - val_loss: 0.0455 - val_accuracy: 0.9867
Epoch 27/50
1050/1050 [=====] - 0s 40us/step - loss: 0.0643 - ac
curacy: 0.9800 - val_loss: 0.0455 - val_accuracy: 0.9844
Epoch 28/50
1050/1050 [=====] - ETA: 0s - loss: 0.0536 - accurac
y: 0.98 - 0s 40us/step - loss: 0.0599 - accuracy: 0.9819 - val_loss: 0.0475 -
val_accuracy: 0.9867
Epoch 00028: early stopping

```

```

In [16]: # plot training history
from matplotlib import pyplot
pyplot.plot(history.history['loss'], label='train')
pyplot.plot(history.history['val_loss'], label='test')
pyplot.legend()
pyplot.show()

```



## Present the model performance

```
In [17]: print('\n# TEST DATA #\n')
score = model.evaluate(X_test, y_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_test[0:10], y_test[0:10], False)

# TEST DATA #

30/30 [=====] - 0s 33us/step
accuracy: 100.00%

# Generate predictions
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
y=Nicola - prediction=Nicola - match=True
y=Theo - prediction=Theo - match=True
y=Jackson - prediction=Jackson - match=True
```

```
In [18]: print('\n# OTHER SPEAKERS DATA #\n')
score = model.evaluate(X_more_test, y_more_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_more_test[0:10], y_more_test[0:10], False)

# OTHER SPEAKERS DATA #

30/30 [=====] - 0s 33us/step
accuracy: 60.00%


# Generate predictions
y=Ankur - prediction=Ankur - match=True
y=Caroline - prediction=Caroline - match=True
y=Rodolfo - prediction=Caroline - match=False
y=Ankur - prediction=Ankur - match=True
y=Caroline - prediction=Caroline - match=True
y=Rodolfo - prediction=Rodolfo - match=True
y=Ankur - prediction=Ankur - match=True
y=Caroline - prediction=Caroline - match=True
y=Rodolfo - prediction=Caroline - match=False
y=Ankur - prediction=Ankur - match=True
```

```
In [19]: print("Classification Report for Test Data\n")
report(X_test, y_test)

print("Classification Report for Other Speakers\n")
report(X_more_test, y_more_test)
```

Classification Report for Test Data


```
[[10  0  0]
 [ 0 10  0]
 [ 0  0 10]]
```



Classification Report					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	10	
1	1.00	1.00	1.00	10	
2	1.00	1.00	1.00	10	
accuracy			1.00	30	
macro avg	1.00	1.00	1.00	30	
weighted avg	1.00	1.00	1.00	30	

Classification Report for Other Speakers

```
[[ 0  0  0  0]
 [ 0 10  0  0]
 [ 2  1  7  0]
 [ 1  1  7  1]]
```



## Classification Report

	precision	recall	f1-score	support
0	0.00	0.00	0.00	0
3	0.83	1.00	0.91	10
4	0.50	0.70	0.58	10
5	1.00	0.10	0.18	10
accuracy			0.60	30
macro avg	0.58	0.45	0.42	30
weighted avg	0.78	0.60	0.56	30

```
c:\users\erodvas\env\lib\site-packages\sklearn\metrics\classification.py:143
9: UndefinedMetricWarning: Recall and F-score are ill-defined and being set t
o 0.0 in labels with no true samples.
  'recall', 'true', average, warn_for)
```

In [ ]: