

Spoken Digits Recognition

Term project for Machine Learning UoT course

The objective of this Notebook is to predict spoken English digits using Keras model.

This Notebook has three main sections.

1. Extract the features of the WAV files, save them into CSV files and store them into Pandas
2. The model will be trained based on the data set created by <https://github.com/Jakobovski/free-spoken-digit-dataset> (<https://github.com/Jakobovski/free-spoken-digit-dataset>)
3. The model will be re-trained based on the previous data plus the recordings made by Ankor (Indian accent), Caroline (Canadian female child accent) and Rodolfo (Brazilian accent)

Section 1

The output of this section is the CSV files with the data to be handle by the model

```
trainData      : ../data/recordings/train
testData       : ../data/recordings/test
moreTrainData  : ../data/recordings/moreSpeakersTrain
moreTestData   : ../data/recordings/moreSpeakersTest
```

```
In [1]: # If true, the WAV files will be read and their features will be saved in the
        CSV files
        # As this is the most time consuming task, only enable it if you don't have th
        e CSV files yet
        CREATE_CSV_FILES = True
```

```
In [2]: # Defines the names of the CSV files
        TRAIN_CSV_FILE = "train.csv"
        TEST_CSV_FILE = "test.csv"
        MORE_TRAIN_CSV_FILE = "more_train.csv"
        MORE_TEST_CSV_FILE = "more_test.csv"
```

```

In [3]: import matplotlib.pyplot as plt
import numpy as np
from matplotlib import cm
import librosa
import csv
import os

def extractWavFeatures(soundFilesFolder, csvFileName):
    print("The features of the files in the folder "+soundFilesFolder+" will be saved to "+csvFileName)
    header = 'filename chroma_stft rmse spectral_centroid spectral_bandwidth rolloff zero_crossing_rate'
    for i in range(1, 21):
        header += f' mfcc{i}'
    header += ' label'
    header = header.split()
    print('CSV Header: ', header)
    file = open(csvFileName, 'w', newline='')
    writer = csv.writer(file)
    writer.writerow(header)
    genres = '1 2 3 4 5 6 7 8 9 0'.split()
    for filename in os.listdir(soundFilesFolder):
        number = f'{soundFilesFolder}/{filename}'
        y, sr = librosa.load(number, mono=True, duration=30)
        # remove leading and trailing silence
        y, index = librosa.effects.trim(y)
        chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
        rmse = librosa.feature.rms(y=y)
        spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
        spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
        rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
        zcr = librosa.feature.zero_crossing_rate(y)
        mfcc = librosa.feature.mfcc(y=y, sr=sr)
        to_append = f'{filename} {np.mean(chroma_stft)} {np.mean(rmse)} {np.mean(spec_cent)} {np.mean(spec_bw)} {np.mean(rolloff)} {np.mean(zcr)}'
        for e in mfcc:
            to_append += f' {np.mean(e)}'
        writer.writerow(to_append.split())
    file.close()
    print("End of extractWavFeatures")

# comment these lines if you already have train.csv and test.csv files
if (CREATE_CSV_FILES == True):
    extractWavFeatures("../data/recordings/train", TRAIN_CSV_FILE)
    extractWavFeatures("../data/recordings/test", TEST_CSV_FILE)
    extractWavFeatures("../data/recordings/moreSpeakersTrain", MORE_TRAIN_CSV_FILE)
    extractWavFeatures("../data/recordings/moreSpeakersTest", MORE_TEST_CSV_FILE)
    print("CSV files are created")
else:
    print("CSV files creation is skipped")

```

The features of the files in the folder ../data/recordings/train will be saved to train.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/test will be saved to test.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/moreSpeakersTrain will be saved to more_train.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

The features of the files in the folder ../data/recordings/moreSpeakersTest will be saved to more_test.csv

```
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid', 'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
```

End of extractWavFeatures

CSV files are created

In [4]: *#Reading a dataset and convert file name to corresponding number*

```
import pandas as pd
import csv
from sklearn import preprocessing

def preProcessData(csvFileName):
    print(csvFileName+ " will be preprocessed")
    data = pd.read_csv(csvFileName)
    data['number'] = data['filename'].str[:1]
    #Dropping unnecessary columns
    data = data.drop(['filename'],axis=1)
    data = data.drop(['label'],axis=1)
    data = data.drop(['chroma_stft'],axis=1)
    data.shape

    print("Preprocessing is finished")
    print(data.head())
    return data

trainData = preProcessData(TRAIN_CSV_FILE)
testData = preProcessData(TEST_CSV_FILE)
moreTrainData = preProcessData(MORE_TRAIN_CSV_FILE)
moreTestData = preProcessData(MORE_TEST_CSV_FILE)
```

train.csv will be preprocessed

Preprocessing is finished

| | rmse | spectral_centroid | spectral_bandwidth | rolloff | \ |
|---|----------|-------------------|--------------------|-------------|---|
| 0 | 0.112672 | 741.829081 | 758.492178 | 1438.494873 | |
| 1 | 0.090344 | 635.610880 | 670.336296 | 1160.452403 | |
| 2 | 0.091456 | 667.786694 | 732.606545 | 1257.180176 | |
| 3 | 0.087751 | 712.304185 | 731.292437 | 1449.104818 | |
| 4 | 0.096603 | 844.363886 | 777.868127 | 1569.583263 | |

| | zero_crossing_rate | mfcc1 | mfcc2 | mfcc3 | mfcc4 | \ |
|---|--------------------|-------------|------------|------------|-----------|---|
| 0 | 0.034023 | -295.578461 | 189.853683 | -19.606564 | 6.078507 | |
| 1 | 0.033458 | -339.148743 | 204.005249 | -7.485526 | 14.297898 | |
| 2 | 0.033268 | -327.507416 | 195.596924 | -3.994768 | 21.315840 | |
| 3 | 0.035916 | -320.809937 | 200.023743 | -8.186146 | 12.661074 | |
| 4 | 0.049465 | -315.801300 | 195.674118 | -13.324564 | 3.544238 | |

| | mfcc5 | ... | mfcc12 | mfcc13 | mfcc14 | mfcc15 | mfcc16 | \ |
|---|-----------|-----|------------|-----------|------------|------------|-----------|---|
| 0 | 22.067095 | ... | -25.725817 | -5.172223 | -8.323026 | -10.299589 | -0.144793 | |
| 1 | 20.885136 | ... | -23.196365 | -1.290891 | -5.515564 | -15.416287 | 0.405876 | |
| 2 | 18.372593 | ... | -18.677113 | -3.098450 | -10.447586 | -10.053793 | 3.248016 | |
| 3 | 15.654718 | ... | -20.832333 | -1.118007 | -6.681235 | -11.685319 | 2.010999 | |
| 4 | 12.279986 | ... | -18.158249 | 6.031695 | -6.353736 | -15.983871 | 1.465030 | |

| | mfcc17 | mfcc18 | mfcc19 | mfcc20 | number |
|---|------------|------------|-----------|------------|--------|
| 0 | -9.017329 | -4.569392 | 2.881349 | -15.627436 | 0 |
| 1 | -3.624587 | -11.204143 | -0.096359 | -6.751650 | 0 |
| 2 | -11.686995 | -10.726046 | 6.857377 | -9.067446 | 0 |
| 3 | -5.946658 | -6.905020 | 4.136240 | -9.614882 | 0 |
| 4 | -5.109472 | -8.666434 | 5.026890 | -5.346444 | 0 |

[5 rows x 26 columns]

test.csv will be preprocessed

Preprocessing is finished

| | rmse | spectral_centroid | spectral_bandwidth | rolloff | \ |
|---|----------|-------------------|--------------------|-------------|---|
| 0 | 0.095394 | 756.450712 | 761.875940 | 1463.941148 | |
| 1 | 0.040176 | 791.046914 | 1039.695939 | 2027.709961 | |
| 2 | 0.006984 | 958.934867 | 941.639039 | 2084.106445 | |
| 3 | 0.071547 | 759.877794 | 899.957003 | 1427.553489 | |
| 4 | 0.030382 | 968.793389 | 1024.834851 | 1911.968994 | |

| | zero_crossing_rate | mfcc1 | mfcc2 | mfcc3 | mfcc4 | \ |
|---|--------------------|-------------|------------|------------|-----------|---|
| 0 | 0.037296 | -328.263885 | 180.479416 | -0.485355 | 15.525293 | |
| 1 | 0.031440 | -385.602570 | 189.328186 | -37.268154 | 59.937920 | |
| 2 | 0.040039 | -542.812622 | 217.971329 | -62.197266 | 21.537390 | |
| 3 | 0.030429 | -355.530396 | 204.388977 | -20.676432 | 26.671131 | |
| 4 | 0.045654 | -376.499390 | 237.137833 | -59.964413 | 37.715607 | |

| | mfcc5 | ... | mfcc12 | mfcc13 | mfcc14 | mfcc15 | mfcc16 | \ |
|---|-----------|-----|------------|-----------|------------|------------|-----------|---|
| 0 | 20.992447 | ... | -15.426966 | 7.284101 | -6.443027 | -13.377846 | -2.407696 | |
| 1 | 45.049831 | ... | -31.051588 | 3.420474 | -9.762264 | -11.220519 | 12.306476 | |
| 2 | 37.756233 | ... | -26.797358 | 5.341060 | -12.159102 | -14.180812 | 9.346475 | |
| 3 | 15.797892 | ... | -18.198524 | 4.029843 | -10.552087 | -21.039103 | -5.634320 | |
| 4 | 21.510382 | ... | -21.805593 | 13.740063 | -6.738161 | -9.305484 | 13.205662 | |

| | mfcc17 | mfcc18 | mfcc19 | mfcc20 | number |
|---|------------|------------|-----------|------------|--------|
| 0 | -12.902534 | -10.437113 | -1.025342 | -15.457672 | 0 |
| 1 | -5.082399 | -3.775387 | 9.707520 | -8.757109 | 0 |

```

2 -10.899978 -9.715154 5.997578 -12.574761 0
3 -14.788965 -11.016036 -1.313916 -16.993853 1
4 -11.917943 -7.877903 9.777577 -10.397771 1

```

[5 rows x 26 columns]

more_train.csv will be preprocessed

Preprocessing is finished

```

      rmse spectral_centroid spectral_bandwidth rolloff \
0 0.039759      1358.208628      1890.243941 2540.917969
1 0.302424      879.994019      1137.986581 1826.501859
2 0.026959      1237.544903      1219.890113 2372.379244
3 0.027274      1523.814892      2030.693021 3053.946533
4 0.304633      625.579402      799.806332 1029.825439

      zero_crossing_rate      mfcc1      mfcc2      mfcc3      mfcc4 \
0      0.093363 -349.631744 135.204880 19.397516 12.044560
1      0.031423 -225.136642 150.581146 -11.930015 5.277394
2      0.059871 -378.096527 180.817047 -41.006123 22.502394
3      0.084253 -403.760406 127.670433 23.293980 7.495267
4      0.023584 -228.231903 191.904144 1.268919 17.871376

      mfcc5 ...      mfcc12      mfcc13      mfcc14      mfcc15      mfcc16 \
0 -8.521679 ... -12.872719 -3.118359 -4.095297 -8.339793 -4.422189
1 -16.297689 ... -13.114679 -13.413434 -17.883400 -17.694012 -18.091360
2 -8.963315 ... -32.601864 -4.715013 -8.889856 -11.225335 -10.951420
3 -6.774144 ... -6.803737 -3.132436 -2.999972 -14.086405 -9.653265
4 -4.571987 ... -5.878148 -9.646009 -15.565687 -17.903820 -14.293053

      mfcc17      mfcc18      mfcc19      mfcc20      number
0 -2.988979 -0.864654 -4.008632 3.243911 0
1 -14.684992 -16.672678 -10.986701 -10.445865 0
2 -8.972343 -4.716431 -5.655877 -7.133625 0
3 -7.696736 -2.747038 -11.928693 -4.345026 1
4 -14.907639 -9.752651 -10.333107 -7.938378 1

```

[5 rows x 26 columns]

more_test.csv will be preprocessed

Preprocessing is finished

```

      rmse spectral_centroid spectral_bandwidth rolloff \
0 0.070496      1736.761057      1697.429353 2782.932447
1 0.178100      986.240750      1243.094994 2207.479581
2 0.159381      1140.769641      1151.467042 2283.347731
3 0.102912      1270.788895      1624.572386 2442.362154
4 0.243922      839.268659      985.135107 1423.724724

      zero_crossing_rate      mfcc1      mfcc2      mfcc3      mfcc4 \
0      0.095979 -270.462311 120.525047 -16.563780 25.340765
1      0.045351 -230.533142 171.157135 -5.515658 20.675375
2      0.056547 -261.164734 169.949051 -12.116495 0.249355
3      0.055965 -264.190643 154.578873 -9.901300 1.276820
4      0.037626 -191.595886 189.683578 -29.157555 16.972124

      mfcc5 ...      mfcc12      mfcc13      mfcc14      mfcc15      mfcc16 \
0 -8.705835 ... -20.140141 -1.016966 -11.883263 -8.389129 -4.095136
1 -3.108402 ... -14.097899 -9.547957 -9.970568 -15.526528 -17.028362
2 8.262510 ... -26.140615 -3.177802 -5.860640 -14.931437 -12.212286
3 -4.987144 ... -9.880266 -0.010168 -10.896443 -5.682155 -0.433506

```

```
4 -13.665516 ... -14.553379 -6.711682 -15.045198 -11.206120 -11.227142
```

| | mfcc17 | mfcc18 | mfcc19 | mfcc20 | number |
|---|------------|------------|------------|------------|--------|
| 0 | 2.231521 | -6.269879 | -7.276823 | 0.429325 | 0 |
| 1 | -13.014733 | -9.094010 | 1.289093 | -5.188507 | 0 |
| 2 | -5.934423 | -2.185390 | -7.319934 | -12.082070 | 0 |
| 3 | 6.786911 | -10.947194 | -11.484864 | -0.297753 | 1 |
| 4 | -9.380514 | -11.611271 | -7.099937 | -6.555156 | 1 |

```
[5 rows x 26 columns]
```

Section 2

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings)
 Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings
 for each digit for each speaker: 30 recordings total. The data used here comes from the recordings stored in:

- ../data/recordings/train
- ../data/recordings/test

```
In [5]: # Splitting the dataset into training, validation and testing dataset
from sklearn.model_selection import train_test_split
X = np.array(trainData.iloc[:, :-1], dtype = float)
y = trainData.iloc[:, -1]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_
state=42)

X_test = np.array(testData.iloc[:, :-1], dtype = float)
y_test = testData.iloc[:, -1]

print("Y from training data:", y_train.shape)
print("Y from validation data:", y_val.shape)
print("Y from test data:", y_test.shape)

Y from training data: (1029,)
Y from validation data: (441,)
Y from test data: (30,)
```

```
In [6]: #Normalizing the dataset
from sklearn.preprocessing import StandardScaler
import numpy as np
scaler = StandardScaler()
X_train = scaler.fit_transform( X_train )
X_val = scaler.transform( X_val )
X_test = scaler.transform( X_test )

print("X from training data", X_train.shape)
print("X from validation data", X_val.shape)
print("X from test data", X_test.shape)
```

```
X from training data (1029, 25)
X from validation data (441, 25)
X from test data (30, 25)
```



```
In [7]: #Creating a Model
from keras import models
from keras import layers
import keras

# model 1
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_shape=(X_train.shape[1],)))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# Learning Process of a model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# simple early stopping
from keras.callbacks import EarlyStopping

es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)

#Train with early stopping to avoid overfitting
history = model.fit(X_train,
                    y_train,
                    validation_data=(X_val, y_val),
                    epochs=50,
                    batch_size=128,
                    callbacks=[es])
```

Using TensorFlow backend.

Train on 1029 samples, validate on 441 samples

Epoch 1/50

1029/1029 [=====] - 1s 486us/step - loss: 2.3728 - accuracy: 0.1409 - val_loss: 2.0880 - val_accuracy: 0.4626

Epoch 2/50

1029/1029 [=====] - 0s 55us/step - loss: 2.1748 - accuracy: 0.1944 - val_loss: 1.9439 - val_accuracy: 0.5533

Epoch 3/50

1029/1029 [=====] - 0s 52us/step - loss: 1.9924 - accuracy: 0.2867 - val_loss: 1.7869 - val_accuracy: 0.6077

Epoch 4/50

1029/1029 [=====] - 0s 47us/step - loss: 1.8279 - accuracy: 0.3615 - val_loss: 1.5884 - val_accuracy: 0.6168

Epoch 5/50

1029/1029 [=====] - 0s 49us/step - loss: 1.7172 - accuracy: 0.4062 - val_loss: 1.3784 - val_accuracy: 0.6327

Epoch 6/50

1029/1029 [=====] - 0s 52us/step - loss: 1.5858 - accuracy: 0.4325 - val_loss: 1.1987 - val_accuracy: 0.7029

Epoch 7/50

1029/1029 [=====] - 0s 52us/step - loss: 1.3984 - accuracy: 0.4956 - val_loss: 1.0512 - val_accuracy: 0.7370

Epoch 8/50

1029/1029 [=====] - 0s 53us/step - loss: 1.3392 - accuracy: 0.5170 - val_loss: 0.9289 - val_accuracy: 0.7619

Epoch 9/50

1029/1029 [=====] - 0s 53us/step - loss: 1.2463 - accuracy: 0.5384 - val_loss: 0.8314 - val_accuracy: 0.7982

Epoch 10/50

1029/1029 [=====] - 0s 56us/step - loss: 1.1302 - accuracy: 0.5870 - val_loss: 0.7613 - val_accuracy: 0.8209

Epoch 11/50

1029/1029 [=====] - 0s 53us/step - loss: 1.0923 - accuracy: 0.6122 - val_loss: 0.6935 - val_accuracy: 0.8367

Epoch 12/50

1029/1029 [=====] - 0s 55us/step - loss: 1.0024 - accuracy: 0.6550 - val_loss: 0.6335 - val_accuracy: 0.8617

Epoch 13/50

1029/1029 [=====] - 0s 53us/step - loss: 0.9564 - accuracy: 0.6569 - val_loss: 0.5766 - val_accuracy: 0.8821

Epoch 14/50

1029/1029 [=====] - 0s 52us/step - loss: 0.8330 - accuracy: 0.7094 - val_loss: 0.5290 - val_accuracy: 0.8844

Epoch 15/50

1029/1029 [=====] - 0s 52us/step - loss: 0.8783 - accuracy: 0.6842 - val_loss: 0.4830 - val_accuracy: 0.8821

Epoch 16/50

1029/1029 [=====] - 0s 54us/step - loss: 0.7899 - accuracy: 0.7153 - val_loss: 0.4461 - val_accuracy: 0.8798

Epoch 17/50

1029/1029 [=====] - 0s 50us/step - loss: 0.8043 - accuracy: 0.7182 - val_loss: 0.4122 - val_accuracy: 0.8889

Epoch 18/50

1029/1029 [=====] - 0s 50us/step - loss: 0.7718 - accuracy: 0.7366 - val_loss: 0.3889 - val_accuracy: 0.8912


Epoch 19/50

1029/1029 [=====] - 0s 50us/step - loss: 0.7500 - ac

```
curacy: 0.7444 - val_loss: 0.3613 - val_accuracy: 0.8980
Epoch 20/50
1029/1029 [=====] - 0s 56us/step - loss: 0.6669 - ac
curacy: 0.7804 - val_loss: 0.3391 - val_accuracy: 0.9048
Epoch 21/50
1029/1029 [=====] - 0s 51us/step - loss: 0.6834 - ac
curacy: 0.7590 - val_loss: 0.3176 - val_accuracy: 0.9206
Epoch 22/50
1029/1029 [=====] - 0s 53us/step - loss: 0.6317 - ac
curacy: 0.7872 - val_loss: 0.2976 - val_accuracy: 0.9365
Epoch 23/50
1029/1029 [=====] - 0s 52us/step - loss: 0.6460 - ac
curacy: 0.7726 - val_loss: 0.2860 - val_accuracy: 0.9252
Epoch 24/50
1029/1029 [=====] - 0s 51us/step - loss: 0.6164 - ac
curacy: 0.7862 - val_loss: 0.2760 - val_accuracy: 0.9274
Epoch 25/50
1029/1029 [=====] - 0s 52us/step - loss: 0.5966 - ac
curacy: 0.7911 - val_loss: 0.2678 - val_accuracy: 0.9297
Epoch 26/50
1029/1029 [=====] - 0s 52us/step - loss: 0.5624 - ac
curacy: 0.8124 - val_loss: 0.2617 - val_accuracy: 0.9252
Epoch 27/50
1029/1029 [=====] - 0s 57us/step - loss: 0.5705 - ac
curacy: 0.8047 - val_loss: 0.2520 - val_accuracy: 0.9388
Epoch 28/50
1029/1029 [=====] - 0s 54us/step - loss: 0.5040 - ac
curacy: 0.8134 - val_loss: 0.2359 - val_accuracy: 0.9433
Epoch 29/50
1029/1029 [=====] - 0s 66us/step - loss: 0.5415 - ac
curacy: 0.8154 - val_loss: 0.2340 - val_accuracy: 0.9456
Epoch 30/50
1029/1029 [=====] - 0s 56us/step - loss: 0.5089 - ac
curacy: 0.8231 - val_loss: 0.2274 - val_accuracy: 0.9456
Epoch 31/50
1029/1029 [=====] - 0s 55us/step - loss: 0.5037 - ac
curacy: 0.8416 - val_loss: 0.2210 - val_accuracy: 0.9478
Epoch 32/50
1029/1029 [=====] - 0s 54us/step - loss: 0.4783 - ac
curacy: 0.8338 - val_loss: 0.2145 - val_accuracy: 0.9501
Epoch 33/50
1029/1029 [=====] - 0s 54us/step - loss: 0.4761 - ac
curacy: 0.8455 - val_loss: 0.2055 - val_accuracy: 0.9501
Epoch 34/50
1029/1029 [=====] - 0s 52us/step - loss: 0.4715 - ac
curacy: 0.8435 - val_loss: 0.1943 - val_accuracy: 0.9546
Epoch 35/50
1029/1029 [=====] - 0s 52us/step - loss: 0.4793 - ac
curacy: 0.8377 - val_loss: 0.1842 - val_accuracy: 0.9592
Epoch 36/50
1029/1029 [=====] - 0s 54us/step - loss: 0.4386 - ac
curacy: 0.8513 - val_loss: 0.1797 - val_accuracy: 0.9569
Epoch 37/50
1029/1029 [=====] - 0s 52us/step - loss: 0.4467 - ac
curacy: 0.8571 - val_loss: 0.1738 - val_accuracy: 0.9615
Epoch 38/50
1029/1029 [=====] - 0s 57us/step - loss: 0.4315 - ac
```

curacy: 0.8533 - val_loss: 0.1782 - val_accuracy: 0.9546
Epoch 00038: early stopping

```
In [8]: # plot training history
from matplotlib import pyplot
pyplot.plot(history.history['loss'], label='train')
pyplot.plot(history.history['val_loss'], label='test')
pyplot.legend()
pyplot.show()
```



Auxiliary functions to show the results

```
In [9]: def printPrediction(X_data, y_data):
        print('\n# Generate predictions')
        for i in range(len(y_data)):
            prediction = model.predict_classes(X_data[i:i+1])
            print("y={}, prediction={}, match={}".format(y_data[i], prediction, y_
data[i]==str(prediction[0])))
```

```
In [10]: import numpy as np
from keras import backend as K
from keras.models import Sequential
from keras.layers.core import Dense, Dropout, Activation, Flatten
from keras.layers.convolutional import Convolution2D, MaxPooling2D
from keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import classification_report, confusion_matrix

def report(X_data, y_data):
    #Confution Matrix and Classification Report
    Y_pred = model.predict_classes(X_data)
    y_test_num = y_data.astype(np.int64)
    print('Confusion Matrix')
    conf_mt = confusion_matrix(y_test_num, Y_pred)
    print(conf_mt)
    plt.matshow(conf_mt)
    plt.show()

    print('\nClassification Report')
    target_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
    print(classification_report(y_test_num, Y_pred, target_names=target_names
    ))
```

Present the model performance

```
In [11]: print('\n# TEST DATA #\n')
score = model.evaluate(X_test, y_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_test[0:10], y_test[0:10])

# TEST DATA #

30/30 [=====] - 0s 33us/step
accuracy: 90.00%


# Generate predictions
y=0, prediction=[0], match=True
y=0, prediction=[3], match=False
y=0, prediction=[0], match=True
y=1, prediction=[1], match=True
y=1, prediction=[4], match=False
y=1, prediction=[1], match=True
y=2, prediction=[2], match=True
y=2, prediction=[2], match=True
y=2, prediction=[3], match=False
y=3, prediction=[3], match=True
```

```
In [12]: print("Classification Report for Test Data\n")
report(X_test, y_test)
```

Classification Report for Test Data

Confusion Matrix

```
[[2 0 0 1 0 0 0 0 0 0]
 [0 2 0 0 1 0 0 0 0 0]
 [0 0 2 1 0 0 0 0 0 0]
 [0 0 0 3 0 0 0 0 0 0]
 [0 0 0 0 3 0 0 0 0 0]
 [0 0 0 0 0 3 0 0 0 0]
 [0 0 0 0 0 0 3 0 0 0]
 [0 0 0 0 0 0 0 3 0 0]
 [0 0 0 0 0 0 0 0 3 0]
 [0 0 0 0 0 0 0 0 0 3]]
```



| Classification Report | | precision | recall | f1-score | support |
|-----------------------|---|-----------|--------|----------|---------|
| | 0 | 1.00 | 0.67 | 0.80 | 3 |
| | 1 | 1.00 | 0.67 | 0.80 | 3 |
| | 2 | 1.00 | 0.67 | 0.80 | 3 |
| | 3 | 0.60 | 1.00 | 0.75 | 3 |
| | 4 | 0.75 | 1.00 | 0.86 | 3 |
| | 5 | 1.00 | 1.00 | 1.00 | 3 |
| | 6 | 1.00 | 1.00 | 1.00 | 3 |
| | 7 | 1.00 | 1.00 | 1.00 | 3 |
| | 8 | 1.00 | 1.00 | 1.00 | 3 |
| | 9 | 1.00 | 1.00 | 1.00 | 3 |
| accuracy | | | | 0.90 | 30 |
| macro avg | | 0.93 | 0.90 | 0.90 | 30 |
| weighted avg | | 0.94 | 0.90 | 0.90 | 30 |

```
In [ ]:
```

Section 3

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings)
 Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings for each digit for each speaker: 30 recordings total.

In addition, there are 2 recordings for each digit for each speaker: Ankur, Caroline and Rodolfo (total 60 recordings) This addition training data has 1 recordings for each digit for each speaker: 30 recordings total. This addition test data has 1 recordings for each digit for each speaker: 30 recordings total.

Therefore the full data set has:

- Training: 1500 recordings
- Training: 60 recordings

The data used here comes from the recordings stored in:

- ../data/recordings/train
- ../data/recordings/test
- ../data/recordings/moreSpeakersTrain
- ../data/recordings/moreSpeakersTest

```
In [13]: # Splitting the dataset into training, validation and testing dataset
from sklearn.model_selection import train_test_split

fullTrainData = trainData.append(moreTrainData)

X = np.array(fullTrainData.iloc[:, :-1], dtype = float)
y = fullTrainData.iloc[:, -1]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_
state=42)

X_test = np.array(testData.iloc[:, :-1], dtype = float)
y_test = testData.iloc[:, -1]

X_more_test = np.array(moreTestData.iloc[:, :-1], dtype = float)
y_more_test = moreTestData.iloc[:, -1]

print("Y from training data:", y_train.shape)
print("Y from validation data:", y_val.shape)
print("Y from test data:", y_test.shape)
print("Y from other speakers test data:", y_more_test.shape)

Y from training data: (1050,)
Y from validation data: (450,)
Y from test data: (30,)
Y from other speakers test data: (30,)
```



```
In [14]: #Normalizing the dataset
from sklearn.preprocessing import StandardScaler
import numpy as np
scaler = StandardScaler()
X_train = scaler.fit_transform( X_train )
X_val = scaler.transform( X_val )
X_test = scaler.transform( X_test )
X_more_test = scaler.transform( X_more_test )

print("X from training data", X_train.shape)
print("X from validation data", X_val.shape)
print("X from test data", X_test.shape)
print("X from other speakers test data", X_more_test.shape)
```

```
X from training data (1050, 25)
X from validation data (450, 25)
X from test data (30, 25)
X from other speakers test data (30, 25)
```

```
In [15]: #Creating a Model
from keras import models
from keras import layers
import keras

# model 1
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_shape=(X_train.shape[1],)))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

# Learning Process of a model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# simple early stopping
from keras.callbacks import EarlyStopping

es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)

#Train with early stopping to avoid overfitting
history = model.fit(X_train,
                   y_train,
                   validation_data=(X_val, y_val),
                   epochs=50,
                   batch_size=128,
                   callbacks=[es])
```

Train on 1050 samples, validate on 450 samples

Epoch 1/50

1050/1050 [=====] - 0s 326us/step - loss: 2.3707 - accuracy: 0.1257 - val_loss: 2.1405 - val_accuracy: 0.4111

Epoch 2/50

1050/1050 [=====] - 0s 41us/step - loss: 2.2013 - accuracy: 0.1838 - val_loss: 2.0223 - val_accuracy: 0.5844

Epoch 3/50

1050/1050 [=====] - 0s 43us/step - loss: 2.0552 - accuracy: 0.2676 - val_loss: 1.8628 - val_accuracy: 0.6178

Epoch 4/50

1050/1050 [=====] - 0s 45us/step - loss: 1.9285 - accuracy: 0.3114 - val_loss: 1.6629 - val_accuracy: 0.6311

Epoch 5/50

1050/1050 [=====] - 0s 54us/step - loss: 1.7719 - accuracy: 0.3781 - val_loss: 1.4478 - val_accuracy: 0.6933

Epoch 6/50

1050/1050 [=====] - 0s 54us/step - loss: 1.6474 - accuracy: 0.4324 - val_loss: 1.2435 - val_accuracy: 0.7244

Epoch 7/50

1050/1050 [=====] - 0s 66us/step - loss: 1.4880 - accuracy: 0.4800 - val_loss: 1.0724 - val_accuracy: 0.7644

Epoch 8/50

1050/1050 [=====] - 0s 60us/step - loss: 1.3651 - accuracy: 0.5248 - val_loss: 0.9377 - val_accuracy: 0.8044

Epoch 9/50

1050/1050 [=====] - 0s 49us/step - loss: 1.2548 - accuracy: 0.5610 - val_loss: 0.8313 - val_accuracy: 0.8489

Epoch 10/50

1050/1050 [=====] - 0s 49us/step - loss: 1.1842 - accuracy: 0.5733 - val_loss: 0.7405 - val_accuracy: 0.8622

Epoch 11/50

1050/1050 [=====] - 0s 50us/step - loss: 1.1079 - accuracy: 0.6162 - val_loss: 0.6721 - val_accuracy: 0.8689

Epoch 12/50

1050/1050 [=====] - 0s 44us/step - loss: 1.0489 - accuracy: 0.6238 - val_loss: 0.6078 - val_accuracy: 0.8822

Epoch 13/50

1050/1050 [=====] - 0s 47us/step - loss: 0.9408 - accuracy: 0.6600 - val_loss: 0.5597 - val_accuracy: 0.8756

Epoch 14/50

1050/1050 [=====] - 0s 49us/step - loss: 0.8966 - accuracy: 0.6981 - val_loss: 0.5216 - val_accuracy: 0.8756

Epoch 15/50

1050/1050 [=====] - 0s 47us/step - loss: 0.8071 - accuracy: 0.7219 - val_loss: 0.4817 - val_accuracy: 0.8911

Epoch 16/50

1050/1050 [=====] - 0s 50us/step - loss: 0.8131 - accuracy: 0.7095 - val_loss: 0.4517 - val_accuracy: 0.8956

Epoch 17/50

1050/1050 [=====] - 0s 46us/step - loss: 0.7349 - accuracy: 0.7514 - val_loss: 0.4240 - val_accuracy: 0.9022

Epoch 18/50

1050/1050 [=====] - 0s 50us/step - loss: 0.6957 - accuracy: 0.7448 - val_loss: 0.4075 - val_accuracy: 0.9111

Epoch 19/50

1050/1050 [=====] - 0s 50us/step - loss: 0.7244 - ac

```
curacy: 0.7486 - val_loss: 0.3930 - val_accuracy: 0.9022
Epoch 20/50
1050/1050 [=====] - 0s 50us/step - loss: 0.7269 - ac
curacy: 0.7552 - val_loss: 0.3761 - val_accuracy: 0.9111
Epoch 21/50
1050/1050 [=====] - 0s 48us/step - loss: 0.6501 - ac
curacy: 0.7857 - val_loss: 0.3621 - val_accuracy: 0.9200
Epoch 22/50
1050/1050 [=====] - 0s 46us/step - loss: 0.6428 - ac
curacy: 0.7752 - val_loss: 0.3444 - val_accuracy: 0.9200
Epoch 23/50
1050/1050 [=====] - 0s 43us/step - loss: 0.6261 - ac
curacy: 0.7829 - val_loss: 0.3298 - val_accuracy: 0.9356
Epoch 24/50
1050/1050 [=====] - 0s 44us/step - loss: 0.5998 - ac
curacy: 0.7867 - val_loss: 0.3245 - val_accuracy: 0.9267
Epoch 25/50
1050/1050 [=====] - 0s 42us/step - loss: 0.5772 - ac
curacy: 0.8200 - val_loss: 0.3143 - val_accuracy: 0.9333
Epoch 26/50
1050/1050 [=====] - 0s 46us/step - loss: 0.5375 - ac
curacy: 0.8162 - val_loss: 0.2933 - val_accuracy: 0.9311
Epoch 27/50
1050/1050 [=====] - 0s 48us/step - loss: 0.5148 - ac
curacy: 0.8410 - val_loss: 0.2826 - val_accuracy: 0.9356
Epoch 28/50
1050/1050 [=====] - 0s 44us/step - loss: 0.5181 - ac
curacy: 0.8343 - val_loss: 0.2777 - val_accuracy: 0.9356
Epoch 29/50
1050/1050 [=====] - 0s 46us/step - loss: 0.5244 - ac
curacy: 0.8257 - val_loss: 0.2727 - val_accuracy: 0.9378
Epoch 30/50
1050/1050 [=====] - 0s 46us/step - loss: 0.4591 - ac
curacy: 0.8410 - val_loss: 0.2674 - val_accuracy: 0.9356
Epoch 31/50
1050/1050 [=====] - 0s 46us/step - loss: 0.4727 - ac
curacy: 0.8352 - val_loss: 0.2571 - val_accuracy: 0.9400
Epoch 32/50
1050/1050 [=====] - 0s 45us/step - loss: 0.4370 - ac
curacy: 0.8552 - val_loss: 0.2488 - val_accuracy: 0.9467
Epoch 33/50
1050/1050 [=====] - 0s 47us/step - loss: 0.3888 - ac
curacy: 0.8705 - val_loss: 0.2452 - val_accuracy: 0.9400
Epoch 34/50
1050/1050 [=====] - 0s 46us/step - loss: 0.4068 - ac
curacy: 0.8629 - val_loss: 0.2399 - val_accuracy: 0.9422
Epoch 35/50
1050/1050 [=====] - 0s 47us/step - loss: 0.4044 - ac
curacy: 0.8752 - val_loss: 0.2356 - val_accuracy: 0.9467
Epoch 36/50
1050/1050 [=====] - 0s 49us/step - loss: 0.3661 - ac
curacy: 0.8838 - val_loss: 0.2320 - val_accuracy: 0.9467
Epoch 37/50
1050/1050 [=====] - 0s 52us/step - loss: 0.3988 - ac
curacy: 0.8724 - val_loss: 0.2296 - val_accuracy: 0.9489
Epoch 38/50
1050/1050 [=====] - 0s 54us/step - loss: 0.3697 - ac
```

curacy: 0.8686 - val_loss: 0.2252 - val_accuracy: 0.9511

Epoch 39/50

1050/1050 [=====] - 0s 46us/step - loss: 0.3831 - ac

curacy: 0.8762 - val_loss: 0.2218 - val_accuracy: 0.9511


Epoch 40/50

1050/1050 [=====] - 0s 42us/step - loss: 0.3764 - ac

curacy: 0.8733 - val_loss: 0.2235 - val_accuracy: 0.9511

Epoch 00040: early stopping

```
In [16]: # plot training history
from matplotlib import pyplot
pyplot.plot(history.history['loss'], label='train')
pyplot.plot(history.history['val_loss'], label='test')
pyplot.legend()
pyplot.show()
```



Present the model performance

```
In [17]: print('\n# TEST DATA #\n')
score = model.evaluate(X_test, y_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_test[0:10], y_test[0:10])

# TEST DATA #

30/30 [=====] - 0s 35us/step
accuracy: 93.33%

# Generate predictions
y=0, prediction=[0], match=True
y=0, prediction=[3], match=False
y=0, prediction=[0], match=True
y=1, prediction=[1], match=True
y=1, prediction=[1], match=True
y=1, prediction=[1], match=True
y=2, prediction=[2], match=True
y=2, prediction=[2], match=True
y=2, prediction=[3], match=False
y=3, prediction=[3], match=True
```

```
In [18]: print('\n# OTHER SPEAKERS DATA #\n')
score = model.evaluate(X_more_test, y_more_test)
print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))

# Prediction
printPrediction(X_more_test[0:10], y_more_test[0:10])

# OTHER SPEAKERS DATA #

30/30 [=====] - 0s 67us/step
accuracy: 30.00%

# Generate predictions
y=0, prediction=[8], match=False
y=0, prediction=[2], match=False
y=0, prediction=[8], match=False
y=1, prediction=[4], match=False
y=1, prediction=[0], match=False
y=1, prediction=[1], match=True
y=2, prediction=[0], match=False
y=2, prediction=[8], match=False
y=2, prediction=[6], match=False
y=3, prediction=[3], match=True
```


```
In [19]: print("Classification Report for Test Data\n")
report(X_test, y_test)

print("Classification Report for Other Speakers\n")
report(X_more_test, y_more_test)
```

Classification Report for Test Data

Confusion Matrix

```
[[2 0 0 1 0 0 0 0 0 0]
 [0 3 0 0 0 0 0 0 0 0]
 [0 0 2 1 0 0 0 0 0 0]
 [0 0 0 3 0 0 0 0 0 0]
 [0 0 0 0 3 0 0 0 0 0]
 [0 0 0 0 0 3 0 0 0 0]
 [0 0 0 0 0 0 3 0 0 0]
 [0 0 0 0 0 0 0 3 0 0]
 [0 0 0 0 0 0 0 0 3 0]
 [0 0 0 0 0 0 0 0 0 3]]
```




| Classification Report | | | | | |
|-----------------------|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 1.00 | 0.67 | 0.80 | 3 | |
| 1 | 1.00 | 1.00 | 1.00 | 3 | |
| 2 | 1.00 | 0.67 | 0.80 | 3 | |
| 3 | 0.60 | 1.00 | 0.75 | 3 | |
| 4 | 1.00 | 1.00 | 1.00 | 3 | |
| 5 | 1.00 | 1.00 | 1.00 | 3 | |
| 6 | 1.00 | 1.00 | 1.00 | 3 | |
| 7 | 1.00 | 1.00 | 1.00 | 3 | |
| 8 | 1.00 | 1.00 | 1.00 | 3 | |
| 9 | 1.00 | 1.00 | 1.00 | 3 | |
| accuracy | | | 0.93 | 30 | |
| macro avg | 0.96 | 0.93 | 0.93 | 30 | |
| weighted avg | 0.96 | 0.93 | 0.94 | 30 | |

Classification Report for Other Speakers

Confusion Matrix

```
[[0 0 1 0 0 0 0 0 2 0]
 [1 1 0 0 1 0 0 0 0 0]
 [1 0 0 0 0 0 1 0 1 0]
 [0 0 0 1 0 0 1 0 1 0]
 [0 1 0 0 1 0 1 0 0 0]
 [0 0 1 0 0 1 1 0 0 0]
 [0 0 0 0 0 0 2 0 1 0]
 [1 0 0 0 0 0 0 0 2 0]
 [0 0 0 0 0 0 1 0 2 0]
 [2 0 0 0 0 0 0 0 0 1]]
```



Classification Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.00 | 0.00 | 0.00 | 3 |
| 1 | 0.50 | 0.33 | 0.40 | 3 |
| 2 | 0.00 | 0.00 | 0.00 | 3 |
| 3 | 1.00 | 0.33 | 0.50 | 3 |
| 4 | 0.50 | 0.33 | 0.40 | 3 |
| 5 | 1.00 | 0.33 | 0.50 | 3 |
| 6 | 0.29 | 0.67 | 0.40 | 3 |
| 7 | 0.00 | 0.00 | 0.00 | 3 |
| 8 | 0.22 | 0.67 | 0.33 | 3 |
| 9 | 1.00 | 0.33 | 0.50 | 3 |
| accuracy | | | 0.30 | 30 |
| macro avg | 0.45 | 0.30 | 0.30 | 30 |
| weighted avg | 0.45 | 0.30 | 0.30 | 30 |

```
c:\users\erodvas\env\lib\site-packages\sklearn\metrics\classification.py:143
7: UndefinedMetricWarning: Precision and F-score are ill-defined and being set
to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)
```

In []: